

## 6.S081 2020 Lecture 4: Virtual Memory

==

## \* plan:

address spaces  
 paging hardware  
 xv6 VM code

## ## Virtual memory overview

## \* today's problem:

[user/kernel diagram]  
 [memory view: diagram with user processes and kernel in memory]  
 suppose the shell has a bug:  
   sometimes it writes to a random memory address  
 how can we keep it from wrecking the kernel?  
   and from wrecking other processes?

## \* we want isolated address spaces

each process has its own memory  
 it can read and write its own memory  
 it cannot read or write anything else  
 challenge:  
   how to multiplex several memories over one physical memory?  
   while maintaining isolation between memories

## \* xv6 uses RISC-V's paging hardware to implement AS's

ask questions! this material is important  
 topic of next lab (and shows up in several other labs)

## \* paging provides a level of indirection for addressing

CPU -> MMU -> RAM  
   VA      PA  
 s/w can only ld/st to virtual addresses, not physical  
 kernel tells MMU how to map each virtual address to a physical address  
 MMU essentially has a table, indexed by va, yielding pa  
   called a "page table"  
   one page table per address space  
 MMU can restrict what virtual addresses user code can use  
 By programming the MMU, the kernel has complete control over va->pa mapping  
 Allows for many interesting OS features/tricks

## \* RISC-V maps 4-KB "pages"

and aligned -- start on 4 KB boundaries  
 4 KB = 12 bits  
 the RISC-V used in xv6 has 64-bit for addresses  
 thus page table index is top 64-12 = 52 bits of VA  
   except that the top 25 of the top 52 are unused  
   no RISC-V has that much memory now  
   can grow in future  
 so, index is 27 bits.

## \* MMU translation

see Figure 3.1 of book  
 use index bits of VA to find a page table entry (PTE)  
 construct physical address using PPN from PTE + offset of VA

## \* what is in PTE?

each PTE is 64 bits, but only 54 are used  
 top 44 bits of PTE are top bits of physical address  
   "physical page number"  
 low 10 bits of PTE flags  
   Present, Writeable, &c  
 note: size virtual addresses != size physical addresses

- \* where is the page table stored?  
in RAM -- MMU loads (and stores) PTEs  
o/s can read/write PTEs  
read/write memory location corresponding to PTEs
- \* would it be reasonable for page table to just be an array of PTEs?  
how big is it?  
 $2^{27}$  is roughly 134 million  
64 bits per entry  
134\*8 MB for a full page table  
wasting roughly 1GB per page table  
one page table per address space  
one address space per application  
would waste lots of memory for small programs!  
you only need mappings for a few hundred pages  
so the rest of the million entries would be there but not needed
- \* RISC-V 64 uses a "three-level page table" to save space  
see figure 3.2 from book  
page directory page (PD)  
PD has 512 PTEs  
PTEs point to another PD or is a leaf  
so  $512*512*512$  PTEs in total  
PD entries can be invalid  
those PTE pages need not exist  
so a page table for a small address space can be small
- \* how does the mmu know where the page table is located in RAM?  
satp holds phys address of top PD  
pages can be anywhere in RAM -- need not be contiguous  
rewrite satp when switching to another address space/application
- \* how does RISC-V paging hardware translate a va?  
need to find the right PTE  
satp register points to PA of top/L2 PD  
top 9 bits index L2 PD to get PA of L1 PD  
next 9 bits index L1 PD to get PA of L0 PD  
next 9 bits index L0 PD to get PA of PTE  
PPN from PTE + low-12 from VA
- \* flags in PTE  
V, R, W, X, U  
xv6 uses all of them
- \* what if V bit not set? or store and W bit not set?  
"page fault"  
forces transfer to kernel  
trap.c in xv6 source  
kernel can just produce error, kill process  
in xv6: "usertrap(): unexpected scause ... pid=... sepc=... stval=..."  
or kernel can install a PTE, resume the process  
e.g. after loading the page of memory from disk
- \* indirection allows paging h/w to solve many problems  
e.g. phys memory doesn't have to be contiguous  
avoids fragmentation  
e.g. lazy allocation (a lab)  
e.g. copy-on-write fork (another lab)  
many more techniques  
topic of next lecture
- \* Q: why use virtual memory in kernel?  
it is clearly good to have page tables for user processes  
but why have a page table for the kernel?

could the kernel run with using only physical addresses?

top-level answer: yes

most standard kernels do use virtual addresses

why do standard kernels do so?

some reasons are lame, some are better, none are fundamental

- the hardware makes it difficult to turn it off

e.g. on entering a system call, one would have to disable VM

- the kernel itself can benefit from virtual addresses

mark text pages X, but data not (helps tracking down bugs)

unmap a page below kernel stack (helps tracking down bugs)

map a page both in user and kernel (helps user/kernel transition)

## ## Virtual memory in xv6

### \* kernel page table

See figure 3.3 of book

simple mapping mostly

map virtual to physical one-on-one

note double-mapping of trampoline

note permissions

why map devices?

### \* each process has its own address space

and its own page table

see figure 3.4 of book

note: trampoline and trapframe aren't writable by user process

kernel switches page tables (i.e. sets satp) when switching processes

### \* Q: why this address space arrangement?

user virtual addresses start at zero

of course user va 0 maps to different pa for each process

16,777,216 GB for user heap to grow contiguously

but needn't have contiguous phys mem -- no fragmentation problem

both kernel and user map trampoline and trapframe page

eases transition user -> kernel and back

kernel doesn't map user applications

not easy for kernel to r/w user memory

need translate user virtual address to kernel virtual address

good for isolation (see spectre attacks)

easy for kernel to r/w physical memory

pa x mapped at va x

### \* Q: does the kernel have to map all of phys mem into its virtual address space?

## ## Code walk through

### \* setup of kernel address space

kvmmap()

Q: what is address 0x10000000 (256M)

Q: how much address space does 1 L2 entry cover? (1G)

Q: how much address space does 1 L1 entry cover? (2MB)

Q: how much address space does 1 L0 entry cover? (4096)

print kernel page table

Q: what is size of address space? (512G)

Q: how much memory is used to represent it after 1st kvmmap()? (3 pages)

Q: how many entries is CLINT? (16 pages)

Q: how many entries is PLIC? (1024 pages, two level 1 PDs)

Q: how many pages is kernel text (8 pages)

Q: how many pages is kernel total (128M = 64 \* 2MB)

Q: Is trampoline mapped twice? (yes, last entry and direct-mapped, entry [2, 3, 7])

kvmminithart();

Q: after executing w\_satp() why will the next instruction be sfence\_vma()?

### \* mappages() in vm.c

arguments are top PD, va, size, pa, perm

- adds mappings from a range of va's to corresponding pa's
  - rounds b/c some uses pass in non-page-aligned addresses
  - for each page-aligned address in the range
    - call walkpgdir to find address of PTE
    - need the PTE's address (not just content) b/c we want to modify
    - put the desired pa into the PTE
    - mark PTE as valid w/ PTE\_P
- \* walk() in vm.c
  - mimics how the paging h/w finds the PTE for an address
  - PX extracts the 9 bits at Level level
  - &pagetable[PX(level, va)] is the address of the relevant PTE
  - if PTE\_V
    - the relevant page-table page already exists
    - PTE2PA extracts the PPN from the PDE
  - if not PTE\_V
    - alloc a page-table page
    - fill in pte with PPN (using PA2PTE)
  - now the PTE we want is in the page-table page
- \* procinit() in proc.c
  - alloc a page for each kernel stack with a guard page
- \* setup user address space
  - allocproc(): allocates empty top-level page table
  - fork(): uvmcopy()
  - exec(): replace proc's page table with a new one
    - uvmmalloc
    - loadseg
  - print user page table for sh
  - Q: what is entry 2?
- \* a process calls sbrk(n) to ask for n more bytes of heap memory
  - user/umalloc.c calls sbrk() to get memory for the allocator
  - each process has a size
    - kernel adds new memory at process's end, increases size
  - sbrk() allocates physical memory (RAM)
  - maps it into the process's page table
  - returns the starting address of the new memory
- \* growproc() in proc.c
  - proc->sz is the process's current size
  - uvmmalloc() does most of the work
  - when switching to user space satp will be loaded with updated page table
- \* uvmmalloc() in vm.c
  - why PGROUNDUP?
  - arguments to mappages()...