

6.S081/6.828 2019 Lecture 11: RISC-V calling convention, stack frames, and gdb

C code is compiled to machine instructions.

How does the machine work at a lower level?

How does this translation work?

How to interact between C and asm

Why this matters: sometimes need to write code not expressible in C

And you need this for the syscall lab!

RISC-V abstract machine

No C-like control flow, no concept of variables, types ...

Base ISA: Program counter, 32 general-purpose registers (x0--x31)

reg	name	saver	description
x0	zero		hardwired zero
x1	ra	caller	return address
x2	sp	callee	stack pointer
x3	gp		global pointer
x4	tp		thread pointer
x5-7	t0-2	caller	temporary registers
x8	s0/fp	callee	saved register / frame pointer
x9	s1	callee	saved register
x10-11	a0-1	caller	function arguments / return values
x12-17	a2-7	caller	function arguments
x18-27	s2-11	callee	saved registers
x28-31	t3-6	caller	temporary registers
pc			program counter

Running example: sum_to(n)

```
int sum_to(int n) {
    int acc = 0;
    for (int i = 0; i <= n; i++) {
        acc += i;
    }
    return acc;
}
```

What does this look like in assembly code?

```
# sum_to(n)
# expects argument in a0
# returns result in a0
sum_to:
    mv t0, a0          # t0 <- a0
    li a0, 0           # a0 <- 0
loop:
    add a0, a0, t0      # a0 <- a0 + t0
    addi t0, t0, -1     # t0 <- t0 - 1
    bnez t0, loop      # if t0 != 0: pc <- loop
    ret
```

Limited abstractions

No typed, positional arguments

No local variables

Only registers

Machine doesn't even see assembly code

Sees binary encoding of machine instructions

Each instruction: 16 bits or 32 bits

E.g. `mv t0, a0` is encoded as 0x82aa

Not quite 1-to-1 encoding from asm, but close

How would another function call `sum_to`?

```
main:
    li a0, 10          # a0 <- 10
    call sum_to
```

What are the semantics of `call`?

```
call label :=
    ra <- pc + 4      ; ra <- address of next instruction
    pc <- label       ; jump to label
```

Machine doesn't understand labels

Translated to either pc-relative or absolute jumps

What are the semantics of `return`?

```
ret :=
    pc <- ra
```

Let's try it out: `demo1.S`

```
(gdb) file user/_demo1
(gdb) break main
(gdb) continue
Why does it stop before running demo1?
(gdb) layout split
(gdb) stepi
(gdb) info registers
(gdb) p $a0
(gdb) advance 18
(gdb) si
(gdb) p $a0
```

What if we wanted a function calling another function?

```
# sum_then_double(n)
# expects argument in a0
# returns result in a0
sum_then_double:
    call sum_to
    li t0, 2          # t0 <- 2
    mul a0, a0, t0    # a0 <- a0 * t0
    ret
```

```
main:
    li a0, 10
    call sum_then_double
```

Let's try it out: `demo2.S`

We get stuck in an infinite loop
Why: overwrote return address (`ra`)

How to fix: save `ra` somewhere

In another register? Won't work, just defers problem.
Solution: save on stack

```
sum_then_double:
    addi sp, sp, 16    # function prologue:
    sd ra, 0(sp)      # make space on stack, save registers
    call sum_to
    li t0, 2
    mul a0, a0, t0
    ld ra, 0(sp)      # function epilogue:
    addi sp, sp, -16  # restore registers, restore stack pointer
    ret
```

Let's try it out: demo3.S

```
(gdb) ...
(gdb) nexti
```

So far, our functions coordinated with each other

This worked because we were writing all the code involved

Could have written it any other way

E.g. passing arguments in t2, getting return value in t3

Conventions surrounding this: "calling convention"

How are arguments passed?

a0, a1, ..., a7, rest on stack

How are values returned?

a0, a1

Who saves registers?

Designated as caller or callee saved

Could ra be a callee-saved register?

Our assembly code should follow this convention

C code generated by GCC follows this convention

This means that everyone's code can interop, incl C/asm interop

Read: demo4.c / demo4.asm

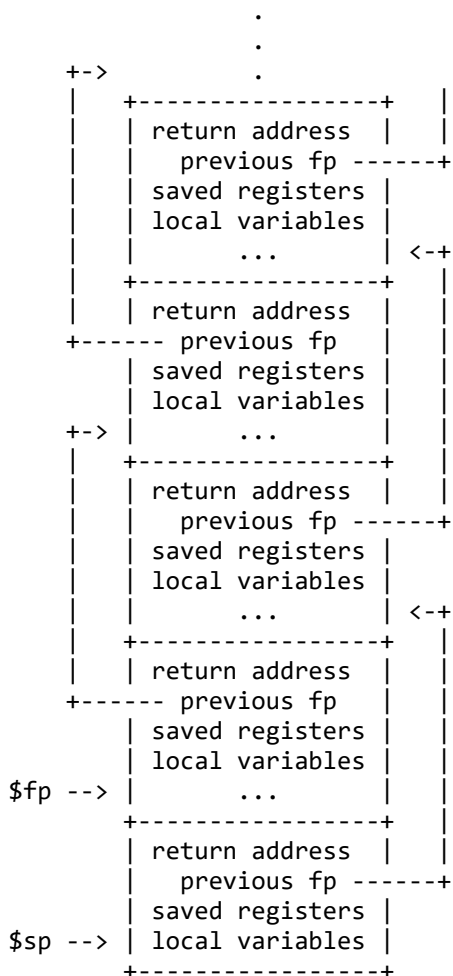
Can see function prologue, body, epilogue

Why doesn't it save ra? Leaf function, not needed

What is going on with s0/fp?

We compiled with -fno-omit-frame-pointer

Stack



Demo program: demo5.c

```
(gdb) break g
```

```
(gdb) si
```

```
(gdb) si
(gdb) si
(gdb) si
(gdb) p $sp
(gdb) p $fp
(gdb) x/g $fp-16
(gdb) x/g 0x0000000000002fd0-16
```

Stack diagram:

```

      0x2fe0 |
      0x2fd8 | <garbage ra>          \
      0x2fd0 | <garbage fp>         / stack frame for main
      0x2fc8 | ra into main        \
$fp --> 0x2fc0 | 0x0000000000002fe0 / stack frame for f
      0x2fb8 | ra into f           \
$sp --> 0x2fb0 | 0x0000000000002fd0 / stack frame for g
```

GDB can automate this reasoning for us

Plus, it can use debug info to reason about leaf functions, etc.

```
(gdb) backtrace
(gdb) info frame
(gdb) frame 1
(gdb) info frame
(gdb) frame 2
(gdb) info frame
```

Calling C from asm / calling asm from C

Follow calling convention and everything will work out

Write function prototype so C knows how to call assembly

Demo: demo6.c / demo6_asm.S

Why do we use s0/s1 instead of e.g. t0/t1?

```
(gdb) b sum_squares_to
(gdb) si ...
(gdb) x/4g $sp
(gdb) si ...
```

Inline assembly

Structs

C struct layout rules

Why: misaligned load/store can be slow or unsupported (platform-dependent)

`__attribute__((packed))`

How to access and manipulate C structs from assembly?

Generally passed by reference

Need to know struct layout

Demo: demo7.c / demo7_asm.S

Debugging

examine: inspect memory contents

```
x/nfu addr
n: count
f: format
u: unit size
```

step/next/finish

step: next line of C code

next: next line of C code, skipping over function calls

finish: continue executing until end of current function call

stepi/nexti

stepi: next assembly instruction

nexti: next assembly instruction, skipping over function calls

layout next

steps through layouts

conditional breakpoints

break, only when a condition holds (e.g. variable has a certain value)

watchpoints

- break when a memory location changes value

GDB is a very powerful tool

- Read the manual for more!

- But you probably don't need all the fancy features for this class

References

RISC-V ISA specification: <https://riscv.org/specifications/>

- Contains detailed information

RISC-V ISA Reference: <https://rv8.io/isa>

- Overview of instructions

RISC-V assembly language reference: <https://rv8.io/asm>

- Overview of directives, pseudo-instructions, and more