

6.S081 2020 Lecture 6: System Call Entry/Exit

Today: user -> kernel transition

system calls, faults, interrupts enter the kernel in the same way
important for isolation and performance
lots of careful design and important detail

What needs to happen when a program makes a system call, e.g. write()?

[CPU | user/kernel diagram]

CPU resources are set up for user execution (not kernel)

32 registers, sp, pc, privilege mode, satp, stvec, sepc, ...

what needs to happen?

save 32 user registers and pc

switch to supervisor mode

switch to kernel page table

switch to kernel stack

jump to kernel C code

high-level goals

don't let user code interfere with user->kernel transition

e.g. don't execute user code in supervisor mode!

transparent to user code -- resume without disturbing

Today we're focusing on the user/kernel transition

and ignoring what the system call implementation does once in the kernel

but the sys call impl has to be careful and secure also!

What does the CPU's "mode" protect?

i.e. what does switching mode from user to supervisor allow?

supervisor can use CPU control registers:

satp -- page table physical address

stvec -- ecall jumps here in kernel; points to trampoline

sepc -- ecall saves user pc here

sscratch -- address of trapframe

supervisor can use PTEs that have no PTE_U flag

but supervisor has no other powers!

e.g. can't use addresses that aren't in page table

so kernel has to carefully set things up so it can work

preview:

write()	write() returns	User
-----	-----	-----
uservec() in trampoline.S	userret() in trampoline.S	Kernel
usertrap() in trap.c	usertrapret() in trap.c	
syscall() in syscall.c	^	
sys_write() in sysfile.c	---	

let's watch an xv6 system call entering/leaving the kernel

xv6 shell writing its \$ prompt

sh.c line 137: write(2, "\$ ", 2);

user/usys.S line 29

this is the write() function, still in user space

a7 tells the kernel what system call we want -- SYS_write = 16

ecall -- triggers the user/kernel transition

let's start by putting a breakpoint on the ecall

user/sh.asm says write()'s ecall is at address 0xde6

\$ make qemu-gdb

(gdb) b *0xde6

(gdb) c

(gdb) delete 1

(gdb) x/3i 0xde4

let's look at the registers

```
(gdb) print $pc
(gdb) info reg
```

\$pc and \$sp are at low addresses -- user memory starts at zero
 C on RISC-V puts function arguments in a0, a1, a2, &c
 write() arguments: a0 is fd, a1 is buf, a2 is n

```
(gdb) x/2c $a1
```

the shell is printing the \$ prompt

what page table is in use?

```
(gdb) print/x $satp
not very useful
qemu: control-a c, info mem
there are mappings for six pages
instructions, data, stack guard (no PTE_U), stack,
then two high mystery pages: trapframe and trampoline
there are no mappings for kernel memory, devices, physical mem
```

let's execute the ecall

```
(gdb) stepi
```

where are we?

```
(gdb) print $pc
we're executing at a very high virtual address
(gdb) x/6i 0x3fffffff000
these are the instructions we're about to execute
see uservec in kernel/trampoline.S
it's the start of the kernel's trap handling code
(gdb) info reg
the registers hold user values (except $pc)
qemu: info mem
we're still using the user page table
note that $pc is in the trampoline page, the very last page
```

we're executing in the "trampoline" page, which contains the start of the kernel's trap handling code. ecall doesn't switch page tables, so these kernel instructions have to exist somewhere in the user page table. the trampoline page is the answer: the kernel maps it at the top of every user page table. the kernel sets \$stvec to the trampoline page's virtual address. the trampoline is protected: no PTE_U flag.

```
(gdb) print/x $stvec
```

can we tell that we're in supervisor mode?

I don't know a way to find the mode directly
 but observe \$pc is executing in a page with no PTE_U flag
 lack of crash implies we are in supervisor mode

how did we get here?

```
ecall did three things:
change mode from user to supervisor
save $pc in $sepc
(gdb) print/x $sepc
jump to $stvec (i.e. set $pc to $stvec)
the kernel previously set $stvec, before jumping to user space
```

note: ecall lets user code switch to supervisor mode
 but the kernel immediately gains control via \$stvec
 so the user program itself can't execute as supervisor

what needs to happen now?

save the 32 user register values (for later transparent resume)

```
switch to kernel page table
set up stack for kernel C code
jump to kernel C code
```

why didn't the RISC-V designers have `ecall` do these things for us?
`ecall` does as little as possible:
 to give O/S designers scope for very fast syscalls / faults / intrs
 maybe O/S can handle some traps w/o switching page tables
 maybe we can map BOTH user and kernel into a single page table
 so no page table switch required
 maybe some registers do not have to be saved
 maybe no stack is required for simple system calls

there have been many clever schemes invented for kernel entry!
 different amounts of work by CPU
 different strategies for handler s/w
 performance here is often super important

what are our options at this point for saving user registers?
 can we just write them somewhere convenient in physical memory?
 no, even supervisor mode is constrained to use the page table
 can we first set `satp` to the kernel page table?
 supervisor mode is allowed to set `satp`...
 but we don't know the address of the kernel page table at this point!
 and we need a free register to even execute `csrw satp, $xx`

two parts to the solution for where to save the 32 user registers:

- 1) xv6 maps a 2nd kernel page, the trapframe, into every user page table
 it has space to hold the saved registers
 the kernel gives each process a different trapframe page
 the page at `0x3fffffe000` is the trapframe page
 see `struct trapframe` in `kernel/proc.h`
 (but we still need a register holding the trapframe's address...)
- 2) RISC-V provides the `sscratch` register
 the kernel puts a pointer to the trapframe in `sscratch`
 before entering user space
 supervisor code can swap any register with `sscratch`
 thus both getting hold of the value in `sscratch`,
 and simultaneously saving the register's user value

see this at the start of `uservec` in `trapframe.S`:
`csrrw a0, sscratch, a0`

the `csrrw` has already been executed due to some gdb quirk...

```
(gdb) print/x $a0
address of the trapframe
(gdb> print/x $sscratch
0x2, the old first argument (fd)
```

now `uservec()` has 32 saves of user registers to the trapframe, via `a0`
 so they can be restored later, when the system call returns
 let's skip them

```
(gdb) b *0x3fffff076
(gdb) c
```

now we're setting up to be able to run C code in the kernel
 first a stack
 previously, kernel put a pointer to top of this process's
 kernel stack in trapframe
 look at `struct trapframe` in `kernel/proc.h`
`"ld sp, 8(a0)"` fetches the kernel stack pointer
 remember `a0` points to the trapframe
 at this point the only kernel data the code can

get at is the trapframe, so everything has to be loaded from there.

```
(gdb) stepi
```

retrieve hart ID into tp

```
(gdb) stepi
```

we want to jump to the kernel C function `usertrap()`, which the kernel previously saved in the trapframe.

"ld t0, 16(a0)" fetches it into t0, we'll use it in a moment, after switching to the kernel page table

```
(gdb) stepi
```

load a pointer to the kernel pagetable from the trapframe, and load it into satp, and issue an sfence to clear the TLB.

```
(gdb) stepi
```

```
(gdb) stepi
```

```
(gdb) stepi
```

why isn't there a crash at this point?

after all we just switched page tables while executing!

answer: the trampoline page is mapped at the same virtual address in the kernel page table as well as every user page table

```
(gdb) print $pc
```

```
qemu: info mem
```

with the kernel page table we can now use kernel functions and data

the `jr t0` is a jump to `usertrap()` (using t0 retrieved from trapframe)

```
(gdb) print/x $t0
```

```
(gdb) x/4i $t0
```

```
(gdb) stepi
```

```
(gdb) tui enable
```

we're now in `usertrap()` in `kernel/trap.c`

various traps come here, e.g. errors, device interrupts, and system calls

`usertrap()` looks in the `scause` register to see the trap cause

see Figure 10.3 on page 102 of The RISC-V Reader

`scause = 8` is a system call

```
(gdb) next ... until syscall()
```

```
(gdb) step
```

```
(gdb) next
```

now we're in `syscall()` `kernel/syscall.c`

`myproc()` uses `tp` to retrieve current struct `proc *`

`p->xxx` is usually a slot in the current process's struct `proc`

`syscall()` retrieves the system call number from saved register `a7`

`p->trapframe` points to the trapframe, with saved registers

`p->trapframe->a7` holds 16, `SYS_write`

`p->trapframe->a0` holds `write()` first argument -- `fd`

`p->trapframe->a1` holds `buf`

`p->trapframe->a2` holds `n`

```
(gdb) next ...
```

```
(gdb) print num
```

then dispatches through `syscall[num]`, a table of functions

```
(gdb) next ...
(gdb) step
```

aha, we're in `sys_write`.
 at this point system call implementations are fairly ordinary C code.
 let's skip to the end, to see how a system call returns to user space.

```
(gdb) finish
```

notice that `write()` produced console output (the shell's \$ prompt)
 back to `syscall()`
 the `p->tf->a0` assignment causes (eventually) `a0` to hold the return value
 the C calling convention on RISC-V puts return values in `a0`

```
(gdb) next
```

```
back to usertrap()
```

```
(gdb) print p->trapframe->a0
```

`write()` returned 2 -- two characters -- \$ and space

```
(gdb) next
(gdb) step
```

now we're in `usertrapret()`, which starts the process of returning
 to the user program

we need to prepare for the next user->kernel transition
`stvec` = `uservec` (the trampoline), for the next `ecall`
`traframe satp` = kernel page table, for next `uservec`
`traframe sp` = top of kernel stack
`trapframe trap` = `usertrap`
`trapframe hartid` = `hartid` (in `tp`)

at the end, we'll use the RISC-V `sret` instruction
 we need to prepare a few registers that `sret` uses
`sstatus` -- set the "previous mode" bit to user
`sepc` -- the saved user program counter (from trap entry)

we're going to switch to the user page table while executing
 not OK in `usertrapret()`, since it's not mapped in the user page table.
 need a page that's mapped in both user and kernel page table -- the trampoline.
 jump to `userret` in trampoline.S

```
(gdb) tui disable
(gdb) step
(gdb) x/8i 0x3fffffff090
```

`a0` holds `TRAPFRAME`
`a1` holds user page table address
 the `csrw satp` switches to the user address space

```
(gdb) stepi
(gdb) stepi
(gdb) stepi
```

the `csrw scratch` puts the user `a0` into `sscratch`
 just before `sret` we'll do a swap,
 so that `a0` holds the user `a0` and `sscratch` holds `trapframe` pointer.
 which is what `uservec` expects.

now 32 loads from the `trapframe` into registers
 these restore the user registers
 let's skip over them

```
(gdb) b *0x3fffffff10a
(gdb) c
```

here's the csrw that swaps a0 with sscratch

```
(gdb) stepi
(gdb) print/x $a0 -- the return value from write()
(gdb) print/x $sscratch -- trapframe address for uservec
```

now we're at the sret instruction

```
(gdb) print $pc
(gdb) stepi
(gdb) print $pc
```

now we're back in the user program (\$pc = 0x0xdea)
returning 2 from the write() function

```
(gdb) print/x $a0
```

and we're done with a system call!

summary

- system call entry/exit is far more complex than function call
- much of the complexity is due to the requirement for isolation
- and the desire for simple and fast hardware mechanisms
- a few design questions to ponder:
 - can an evil program abuse the entry mechanism?
 - can you think of ways to make the hardware or software simpler?
 - can you think of ways to make traps faster?