

6.S081 2020 Lecture 9: Interrupts

Interrupts

hardware wants attention now!

e.g., pkt arrived, clock interrupt

software must set aside current work and respond

on RISC-V use same trap mechanism as for syscalls and exceptions

new issues/complications:

asynchronous

interrupts running process

interrupt handler may not run in context of process who caused interrupt

concurrency

devices and process run in parallel

programming devices

device can be difficult to program

Where do device interrupts come from?

diagram: Fig 1 of SiFive board in FU540-C000-v1.0.pdf

CPUs, CLINT, PLIC, devices

Fig 3 has more detail for interrupts

Section 8 and 9 of FU540-C000-v1.0.pdf

UART (universal asynchronous receiver/transmitter)

See Section 13 of FU540-C000-v1.0.pdf

Although the QEMU version is slightly different

Follows <http://byterunner.com/16550.html>

the interrupt tells the kernel the device hardware wants attention

the driver (in the kernel) knows how to tell the device to do things

often the interrupt handler calls the relevant driver

but other arrangements are possible (schedule a thread; poll)

[diagram: top-half/bottom-half]

Case study: console output and keyboard input

\$ ls

how does \$ show up on console?

printing to simulated console:

driver puts characters into UART's send FIFO

interrupt when character has been sent

informs driver that it can send more

how are the characters l and s read from keyboard (and echo to console)?

reading from simulated keyboard

user hits key, which returns in UART interrupt

driver gets character from UART's receive FIFO

how does kernel know which device interrupted?

each device has a unique source/IRQ (interrupt request) number

defined by hardware platform

UART0 is 10 on qemu (see kernel/memlayout.h)

different on SiFive board

RISC-V interrupt-related registers

sie --- supervisor interrupt enabled register

one bit per software interrupt, external interrupt, timer interrupt

sstatus --- supervisor status register

one bit to enable interrupts

sip --- supervisor interrupt pending register

scause --- supervisor cause register

stvec --- supervisor trap vector register

mdeleg --- machine delegate register

Let's look at how xv6 sets up the interrupt machinery

start():

w_sie(r_sie() | SIE_SEIE | SIE_STIE | SIE_SSIE);

main():

consoleinit();

```

    uartinit()
    plicinit();
    scheduler();
    intr_on();
    w_sstatus(r_sstatus() | SSTATUS_SIE);

```

Printing "\$"

shell is started with fd 0, 1, 2 for "console"
 setup by init
 Unix presents console device as a file!

```

printf()
putc()
    write(2, "$", 1)

```

```

sys_write()
    filewrite()
        consolewrite()
            uartputc()
                add "$" to buffer
            uartstart() // kick off sending character
                put the character in UART fifo
                return to user space ...
                while at the same time UART is sending character to console

```

shell calls `sys_read()` to wait for input (see below)

UART completes sending character and raises interrupt

PLIC passes interrupt on to a CPU

the CPU performs the following steps on interrupt:

- If the trap is a device interrupt, and the SIE bit is clear, don't do any of the following.
 - Disable interrupts by clearing SIE.
Prevents interrupts being interrupted
 - Copy the pc to sepc
 - Save the current mode (user or supervisor) in the SPP bit in sstatus.
 - Set scause to reflect the interrupt's cause.
 - Set the mode to supervisor.
 - Copy stvec to the pc.
stvec contains kernelvec or usertrap, depending if interrupt happened in user space or kernel space
 - Start executing at the new pc.
- same mechanism we have seen before systems calls, pgfaults, etc.

both kernelvec/usertrap call `devintr()` to check for interrupts

external interrupt for UART

```

    plic_claim()
    uartintr()
        if there are more characters send them
        may send multiple characters
    plic_complete()
    return from kernelvec/usertrap
    resumes interrupted computation

```

What if several interrupts arrive?

The PLIC distributes interrupts among cores

Interrupts can be handled in parallel

If no CPU claims the interrupt, the interrupt stays pending

Eventually each interrupt is delivered to some CPU

Interrupts expose several forms of concurrency

1. Between device and CPU
Producer/consumer parallelism
2. Interrupt may interrupt the CPU that is returning to shell (still in kernel)
Disable interrupts when code must be atomic

3. Interrupt may run on different CPU in parallel with shell (or returning to shell)
 Locks; topic for Wed

Producer/consumer parallelism

For printing

shell is producer

device is consumer

To decouple the two:

a buffer in the driver

top-half puts chars into buffer

wait if there is no room

runs in the context of the calling process

bottom half remove chars from buffer

interrupt handler wakes up producers

may not run the context of the shell

Note: bottom half and top half may run in parallel on different CPUs

We will get to this in a later lecture

Interrupts interrupt running code

Interrupts run between my code

For example, my code is

1. addi sp,sp,-48

2. sd ra,40(sp)

Q: Might other code run between 1 and 2?

Yes!

Interrupt may happen between 1 and 2

e.g., timer interrupt or uart interrupt

For user code maybe not that bad

Kernel will resume user code in in the same state

For kernel code could be difficult

Interrupt handler may update state that is observable by my code

my code: interrupt:

x = 0

if x = 0 then x = 1

f()

f() may be executed or may not be executed!

To make a block of code "atomic", turn off interrupts

intr_off(): w_sstatus(r_sstatus() & ~SSTATUS_SIE);

RISC-V turns of interrupt on a trap (interrupt/exception)

Can kernel handle interrupt in trampoline.S?

Our first glimps of "concurrency"

We'll get back to this when discussing locking

\$

shell is in read system call to get input from console

usertrap() for system call

w_stvec((uint64)kernelvec);

consoleread()

sleep()

scheduler()

intr_on()

\$ 1

user hits 1, which causes UART interrupt

kernelvec:

save space on current stack; which stack?

save registers on the current stack

in our example, the scheduler thread's stack

kerneltrap()

devintr()

uartintr()

c = uartgetc()

consoleintr(c)

```

    handle ctrl characters
    echo character ('l') using uartput_sync()
    put c in buffer
    wakeup reader
    return from devintr()
return from kerneltrap()
load registers back
sret
Q: where does sret return too
    where ever the interrupt happened (in scheduler loop in this case)
scheduler runs shell so that it can collect 'l'

```

Producer/consumer parallelism

```

For reading from keyboard opposite from printing
    shell is consumer
    device is producer

```

Interrupt evolution

```

Interrupts used to be relatively fast; now they are slow
    old approach: every event causes an interrupt, simple h/w, smart s/w
    new approach: h/w completes lots of work before interrupting
Some devices generate events faster than one per microsecond
    e.g. gigabit ethernet can deliver 1.5 million small packets / second
An interrupt takes on the order of a microsecond
    save/restore state
    cache misses
    what to do if interrupt comes in faster than 1 per microsecond?

```

Polling: another way of interacting with devices

```

Processor spins until device wants attention
    Wastes processor cycles if device is slow
    One example in xv6: uartputc_sync()
But inexpensive if device is fast
    No saving of registers etc.
If events are always waiting, no need to keep alerting the software

```

Polling versus interrupts

```

Polling rather than interrupting, for high-rate devices
Interrupt for low-rate devices, e.g. keyboard
    constant polling would waste CPU
Switch between polling and interrupting automatically
    interrupt when rate is low (and polling would waste CPU cycles)
    poll when rate is high (and interrupting would waste CPU cycles)
Faster forwarding of interrupts to user space
    for page faults and user-handled devices
    h/w delivers directly to user, w/o kernel intervention?
    faster forwarding path through kernel?
We will be seeing many of these topics later in the course

```