

6.S081 2020 Lecture 10: Locking

Why talk about locking?

- apps want to use multi-core processors for parallel speed-up
- so kernel must deal with parallel system calls
- and thus parallel access to kernel data (buffer cache, processes, &c)
- locks help with correct sharing of data
- locks can limit parallel speedup

What goes wrong if we don't have locks

Case study: delete acquire/release in kalloc.c

Boot

kernel works!

Run usertests

- all tests pass!
- except we lose some pages

Why do we lose pages?

- picture of shared-memory multiprocessor
- race between two cores calling kfree()

BUMMER:

- we need locks for correctness
- but loose performance (kfree is serialized)

The lock abstraction:

```
lock l
acquire(l)
  x = x + 1 -- "critical section"
release(l)
a lock is itself an object
if multiple threads call acquire(l)
  only one will return right away
  the others will wait for release() -- "block"
a program typically has lots of data, lots of locks
if different threads use different data,
  then they likely hold different locks,
  so they can execute in parallel -- get more work done.
note that lock l is not specifically tied to data x
the programmer has a plan for the correspondence
```

A conservative rule to decide when you need to lock:

- any time two threads use a memory location, and at least one is a write
- don't touch shared data unless you hold the right lock!
- (too strict: program logic may sometimes rule out sharing; lock-free)
- (too loose: printf(); not always simple lock/data correspondence)

Could locking be automatic?

- perhaps the language could associate a lock with every data object
- compiler adds acquire/release around every use
- less room for programmer to forget!

that idea is often too rigid:

```
rename("d1/x", "d2/y"):
```

```
  lock d1, erase x, unlock d1
```

```
  lock d2, add y, unlock d2
```

problem: the file didn't exist for a while!

```
  rename() should be atomic
```

- other system calls should see before, or after, not in between
- otherwise too hard to write programs

we need:

```
  lock d1 ; lock d2
```

```
  erase x, add y
```

```
  unlock d2; unlock d1
```

- that is, programmer often needs explicit control over
- the region of code during which a lock is held
- in order to hide awkward intermediate states

Ways to think about what locks achieve

- locks help avoid lost updates
- locks help you create atomic multi-step operations -- hide intermediate states
- locks help operations maintain invariants on a data structure
 - assume the invariants are true at start of operation
 - operation uses locks to hide temporary violation of invariants
 - operation restores invariants before releasing locks

Problem: deadlock

notice rename() held two locks

what if:

core A	core B
rename(d1/x, d2/y)	rename(d2/a, d1/b)
lock d1	lock d2
lock d2 ...	lock d1 ...

solution:

- programmer works out an order for all locks
- all code must acquire locks in that order
- i.e. predict locks, sort, acquire -- complex!

Locks versus modularity

- locks make it hard to hide details inside modules
- to avoid deadlock, I need to know locks acquired by functions I call
- and I may need to acquire them before calling, even if I don't use them
- i.e. locks are often not the private business of individual modules

Locks and parallelism

- locks *prevent* parallel execution
- to get parallelism, you often need to split up data and locks
 - in a way that lets each core use different data and different locks
 - "fine grained locks"
- choosing best split of data/locks is a design challenge
 - whole FS; directory/file; disk block
 - whole kernel; each subsystem; each object
- you may need to re-design code to make it work well in parallel
 - example: break single free memory list into per-core free lists
 - helps if threads were waiting a lot on lock for single free list
 - such re-writes can require a lot of work!

Lock granularity advice

- start with big locks, e.g. one lock protecting entire module
 - less deadlock since less opportunity to hold two locks
 - less reasoning about invariants/atomicity required
- measure to see if there's a problem
 - big locks are often enough -- maybe little time spent in that module
- re-design for fine-grained locking only if you have to

Let's look at locking in xv6.

A typical use of locks: uart.c

typical of many O/S's device driver arrangements

diagram:

user processes, kernel, UART, uartputc, remove from uart_tx_buf,
uartintr()

sources of concurrency: processes, interrupt

only one lock in uart.c: uart_tx_lock -- fairly coarse-grained

uartputc() -- what does uart_tx_lock protect?

1. no races in uart_tx_buf operations
2. if queue not empty, UART h/w is executing head of queue
3. no concurrent access to UART write registers

uartintr() -- interrupt handler

- acquires lock -- might have to wait at interrupt level!
- removes character from uart_tx_buf
- hands next queued char to UART h/w (2)

touches UART h/w registers (3)

How to implement locks?

why not:

```
struct lock { int locked; }
acquire(l) {
    while(1){
        if(l->locked == 0){ // A
            l->locked = 1;    // B
            return;
        }
    }
}
```

oops: race between lines A and B

how can we do A and B atomically?

Atomic swap instruction:

```
a5 = 1
s1 = &lk->locked
amoswap.w.aq a5, a5, (s1)
```

does this in hardware:

```
lock addr globally (other cores cannot use it)
temp = *s1
*addr = a5
a5 = temp
unlock addr
```

RISC-V h/w provides a notion of locking a memory location

different CPUs have had different implementations

diagram: cores, bus, RAM, lock thing

so we are really pushing the problem down to the hardware

h/w implements at granularity of cache-line or entire bus

memory lock forces concurrent swap to run one at a time, not interleaved

Look at xv6 spinlock implementation

```
acquire(l){
    while(__sync_lock_test_and_set(&lk->locked, 1) != 0)
}
if l->locked was already 1, sync_lock_test_and_set sets to 1 (again), returns 1,
and the loop continues to spin
if l->locked was 0, at most one lock_test_and_set will see the 0; it will set
it to 1 and return 0; other test_and_set will return 1
this is a "spin lock", since waiting cores "spin" in acquire loop
```

what is the push_off() about?

why disable interrupts?

release():

sets lk->locked = 0

and re-enables interrupts

Detail: memory read/write ordering

suppose two cores use a lock to guard a counter, x

and we have a naive lock implementation

Core A:	Core B:
locked = 1	
x = x + 1	while(locked == 1)
locked = 0	...
	locked = 1
	x = x + 1
	locked = 0

the compiler AND the CPU re-order memory accesses

i.e. they do not obey the source program's order of memory references

e.g. the compiler might generate this code for core A:

```
locked = 1
```

```
locked = 0
x = x + 1
i.e. move the increment outside the critical section!
the legal behaviors are called the "memory model"
release()'s call to __sync_synchronize() prevents re-order
compiler won't move a memory reference past a __sync_synchronize()
and (may) issue "memory barrier" instruction to tell the CPU
acquire()'s call to __sync_synchronize() has a similar effect:
if you use locks, you don't need to understand the memory ordering rules
you need them if you want to write exotic "lock-free" code
```

Why spin locks?

```
don't they waste CPU while waiting?
why not give up the CPU and switch to another process, let it run?
what if holding thread needs to run; shouldn't waiting thread yield CPU?
spin lock guidelines:
  hold spin locks for very short times
  don't yield CPU while holding a spin lock
systems provide "blocking" locks for longer critical sections
  waiting threads yield the CPU
  but overheads are typically higher
you'll see some xv6 blocking schemes later
```

Advice:

```
don't share if you don't have to
start with a few coarse-grained locks
instrument your code -- which locks are preventing parallelism?
use fine-grained locks only as needed for parallel performance
use an automated race detector
```