

## 6.S081/6.828 2019 Lecture 13: Crash Recovery, Logging

## Plan

problem: crash recovery  
 crash leads to inconsistent on-disk file system  
 solution:  
 logging

This is the last xv6 lecture  
 next week we'll switch to papers

## # Why crash recovery

## What is crash recovery?

you're writing the file system  
 then the power fails  
 you reboot  
 is your file system still useable?

## the problem:

crash during multi-step operation  
 may leave FS invariants violated  
 after reboot:  
 bad: crash again due to corrupt FS  
 worse: no crash, but reads/writes incorrect data

## examples:

## create:

```
$ echo hi > x
// create trace from last lecture:
bwrite: block 33 by ialloc // allocate inode in inode block 33
bwrite: block 33 by iupdate // update inode (e.g., set nlink)
bwrite: block 46 by writei // write directory entry, adding "x" by dirlink()
bwrite: block 32 by iupdate // update directory inode, because inode may have changed
```

crash between iupdate and writei  
 allocate file inode  
 crash: inode not free but not used -- not so bad

what if the file system first wrote 46 and 32 and then 33  
 if crash between 32 and 33, then dirent points to free inode -- disaster!  
 crash again, or worse if inode is allocated for something else

## write:

inode addr[] and len  
 indirect block  
 block content  
 block free bitmap  
 crash: inode refers to free block -- disaster!  
 crash: block not free but not used -- not so bad

## unlink:

block free bitmaps  
 free inode  
 erase dirent

## what can we hope for?

- after rebooting and running recovery code
1. FS internal invariants maintained  
 e.g., no block is both in free list and in a file
  2. all but last few operations preserved on disk  
 e.g., data I wrote yesterday are preserved  
 but perhaps not data I was writing at time of crash

so user might have to check last few operations

3. no order anomalies

echo 99 > result ; echo done > status

correctness and performance often conflict

disk writes are slow!

safety => write to disk ASAP

speed => don't write the disk (batch, write-back cache, sort by track, &c)

crash recovery is a recurring problem

arises in all storage systems, e.g. databases

a lot of work has gone into solutions over the years

many clever performance/correctness tradeoffs

# Logging solution

most popular solution: logging (== journaling)

goal: atomic system calls w.r.t. crashes

goal: fast recovery (no hour-long fsck)

will introduce logging in two steps

first xv6's log, which only provides safety and fast recovery

then Linux EXT3, which is also fast in normal operation

the basic idea behind logging

you want atomicity: all of a system call's writes, or none

let's call an atomic operation a "transaction"

record all writes the sys call \*will\* do in the log on disk (log)

then record "done" on disk (commit)

then do the FS disk writes (install)

on crash+recovery:

if "done" in log, replay all writes in log

if no "done", ignore log

this is a WRITE-AHEAD LOG

write-ahead log rule

install \*none\* of a transaction's writes to disk

until \*all\* writes are in the log on disk,

and the logged writes are marked committed.

why the rule?

once we've installed one write to the on-disk FS,

we have to do \*all\* of the transaction's other

writes -- so the transaction is atomic. we have

to be prepared for a crash after the first installation

write, so the other writes must be still available

after the crash -- in the log.

logging is magic

crash recovery of complex mutable data structures is generally hard

logging can often be layered on existing storage systems

and it's compatible with high performance (topic for next week)

# Overview of xv6 logging

xv6 log representation

[diagram: buffer cache, in-memory log block # array,  
FS tree on disk, log header and blocks on disk]

on write add blockno to in-memory array

keep the data itself in buffer cache (pinned)

on commit:

write buffers to the log on disk

WAIT for disk to complete the writes ("synchronous")

write the log header sector to disk

block numbers

```

    non-zero "n"
after commit:
    install (write) the blocks in the log to their home location in FS
    unpin blocks
    write zero to "n" in the log header on disk

```

the "n" value in the log header on disk indicates commit  
 non-zero == committed, log content valid and is a complete transaction  
 zero == not committed, may not be complete, recovery should ignore log  
 write of non-zero "n" is the "commit point"

xv6 disk layout with block numbers

```

2: log head
3: logged blocks
32: inodes
45: bitmap
46: content blocks

```

Let's look at an example.

I've modified bwrite() to print low-level disk writes,  
 i.e. the disk writes that occur during transaction commit.

```
$ echo a > x
```

```

// create
bwrite 3    // inode, 33
bwrite 4    // directory content, 46
bwrite 5    // directory iode, 32
bwrite 2    // commit (block #s and n)
bwrite 33   // install inode for x
bwrite 46   // install directory content
bwrite 32   // install dir inode
bwrite 2    // mark log "empty"
// write
bwrite 3
bwrite 4
bwrite 5
bwrite 2
bwrite 45   // bitmap
bwrite 595  // a (note: bzero was absorbed)
bwrite 33   // inode (file size)
bwrite 2
// write
bwrite 3
bwrite 4
bwrite 2
bwrite 595  // \n
bwrite 33   // inode
bwrite 2

```

let's look at the 2nd transaction, a write()

```

first file.c:syswrite
    compute how many blocks we can write before log is full
    write that many blocks in a transaction

```

```

combined with fs.c:writei
    begin_op()
        bmap() -- can write bitmap, indirect block
        log_write to bzero new block
    bread()
    modify bp->data
    log_write()
        absorbs bzero
    iupdate() -- writes inode
    end_op()

```

```

begin_op() in log.c:
    need to indicate which group of writes must be atomic!
    need to check if log is being committed
    need to check if our writes will fit in remainder of log
log_write():
    add sector # to in-memory array
    bpin() will pin block in buffer cache, so that bio.c won't evict it
end_op():
    if no outstanding operations, commit
commit():
    copy updated blocks from cache to log on disk
    record sector #s and "done" in on-disk log header
    install writes -- copy from on-disk log to on-disk FS
    bunpin() will unpin from cache --- now it can be evicted
    erase "done" from log

```

What would have happened if we crashed during a transaction?

memory is lost, leaving only the disk as of the crash

kernel calls `recover_from_log()` during boot, before using FS

if log header block says "done":

- copy blocks from log to real locations on disk

what is in the on-disk log?

- crash before commit
- crash during commit -- commit point?
- crash during `install_trans()`
- crash just after reboot, while in `recover_from_log()`

note: it is OK to replay the log more than once!

as long no other activity intervenes

note xv6 assumes the disk is fail-stop

- it either does the write correctly, or does not do the write
- i.e. perhaps it can't do the last write due to power failure

thus:

- no partial writes (each sector write is atomic)
- no wild writes
- no decay of sectors (no read errors)
- no read of the wrong sector

## # Challenges

challenge: prevent write-back from cache

- a system call can safely update a `*cached*` block,
- but the block cannot be written to the FS
- until the transaction commits

tricky because e.g. cache may run out of space,

- and be tempted to evict some entries in order
- to read and cache other data.

consider create example:

- write dirty inode to log
- write dir block to log
- evict dirty inode
- commit

xv6 solution:

- ensure buffer cache is big enough
- pin dirty blocks in buffer cache
- after commit, unpin block

challenge: system's call data must fit in log

xv6 solution:

- compute an upper bound of number of blocks each system call writes
- set log size  $\geq$  upper bound
- break up some system calls into several transactions
- for example, large `write()`s
- thus: large `write()`s are not atomic

but a crash will leave a correct prefix of the write

challenge: allowing concurrent system calls

- must allow writes from several calls to be in log

- on commit must write them all

- BUT cannot write data from calls still in a transaction

xv6 solution

- allow no new system calls to start if their data might not fit in log

- must wait for current calls to complete and commit

- when number of in-progress calls falls to zero

- commit

- free up log space

- wake up waiting calls

challenge: a block may be written multiple times in a transaction

- writes affect only the cached block in memory

- so a cached block may reflect multiple uncommitted transactions

- but install only happens when there are no in-progress transactions

- so installed blocks reflect only committed transactions

- good for performance: "write absorbtion"

# Summary

what is good about xv6's log design?

- correctness due to write-ahead log

- good disk throughput: log naturally batches writes

- but data disk blocks are written twice

- concurrency

what's wrong with xv6's logging?

- not very efficient:

- every block is written twice (log and install)

- logs whole blocks even if only a few bytes modified

- writes each log block synchronously

- could write them as a batch and only write head synchronously

- log writes and install writes are eager

- both could be lazy, for more write absorbtion

- but must still write the log first

- trouble with operations that don't fit in the log

- unlink might dirty many blocks while truncating file