

## 6.S081 2020 Lecture 17: User-level virtual memory

Reading: VM primitives for user programs by Appel & Li (1991)

## Plan:

- OS kernel uses virtual memory in creative ways
- Paper argues user-level application can also benefit from VM
  - Concurrent garbage collector
  - Generation garbage collector
  - Concurrent check-pointing
  - Data-compression paging
  - Persistent stores
- Most OSes have mmap() and user-level pagefaults

## What VM user-level VM primitives?

- Trap: handle page-fault traps in user mode
- Prot1: decrease the accessibility of a page
- ProtN: decrease the accessibility of N pages
- Unprot: increase the accessibility of a page
- Dirty: return a list of dirtied pages since previous calls
- Map2: map the same physical page at two different VAs with different prot levels.

Xv6 supports none of them

One way to see the paper is that a good OS should support them  
What about Unix today?

## Unix today: mmap()

Maps memory into the address space (many flags and options)  
Example: mapping file  
mmap(NULL, len, PROT\_READ|PROT\_WRITE, MAP\_PRIVATE, fd, offset)  
Examples: anonymous memory  
mmap(NULL, len, PROT\_READ|PROT\_WRITE, MAP\_PRIVATE|MAP\_ANONYMOUS, -1, 0)  
(preferred over sbrk)

## Unix today: mprotect()

Changes the permission of a mapping  
Examples:  
mprotect(addr, len, PROT\_READ)  
mprotect(addr, len, PROT\_NONE)  
results in a trap on access

## Unix today: munmap()

Removes a mapping  
Example:  
munmap(addr, len)

## Unix today: sigaction()

Configures a signal handler  
(Think sigalarm() lab)  
act.sa\_sigaction = handle\_sigsegv;  
act.sa\_flags = SA\_SIGINFO;  
sigemptyset(&act.sa\_mask)  
sigaction(SIGSEGV, &act, NULL);

## Additional Linux VM calls

Madvise(), Mincore(), Mremap(), Msync(), Mlock(), Mbind(), Shmat(),...

## VM Implementation

Address spaces consists of VMAs and page tables  
VMA (virtual memory area)  
contiguous range of virtual addresses  
same permissions  
backed by the same object (file, anonymous memory)

VMA helps the kernel to decide how to handle page faults

Trap/sigaction implementation

PTE (or TLB entry) marked "protected"  
 CPU saves user state, jumps into kernel.  
 Kernel asks VM system what to do?  
   I.e. page in from disk? Core dump?  
   VM system looks at VMA  
 Generate signal -- upcall into user process.  
   Lower on user stack, or on separate stack...  
 Run user handler, can do anything.  
   Probably must call UNPROT for referenced page.  
   That is, must avoid repeated fault.  
 User handler returns to kernel.  
 Kernel returns to user program.  
 Continue or re-start instruction that trapped.

Can we support user-level VM primitives?

Trap: sigaction and SIGSEGV  
 Prot1: mprotect()  
 ProtN: mprotect()  
 Unprot: mprotect()  
 Dirty: No, but workaround exists  
 Map2: Not directly, but shm\_open/mmap/mmap

Demo: large sqrt table backed by a single page

Application code thinks there is a pre-computed big table  
   n -> sqrt(n)  
 Application can lookup sqrt: sqrts[n]  
   If present, very fast!  
 Table is bigger than physical memory  
 User-level VM primitives allow it to occupy only a \*single\* page

Use case: Concurrent GC

Application allocates memory and computes with it  
   Application is called the "mutator"  
   Application doesn't have to call free  
   (which is error prone, in particular in concurrent programs)  
 Collector concurrently finds free memory  
   Memory that isn't in use by application anymore  
   Not reachable from root pointer or registers of any thread  
 Traditional implementation requires language/compiler support  
   Instrument loads/stores  
 User-level virtual memory can avoid instrumenting load/stores  
   Faster

Example: Baker's real-time GC algorithm (see toy version in baker.c)

<https://dl.acm.org/citation.cfm?id=359460.359470>  
 Divide heap into 2 regions: from-space and to-space  
   to-space is further divided in: scanned, unscanned, and new  
 At start of collection: all objects are in from-space  
   Copy roots to to-space (register and stack)  
   Put forwarding pointers in from-space to the to-space copy  
 At the end of scan, from-space is free memory  
 Don't have to stop the world for GC  
   Do a little of scanning on each allocation  
   Or when dereferencing pointer in from-space

Observations

Why attractive?  
   Alloc is cheap. Compacts, so no free list.  
   Incremental: every allocation scan/copy a little  
   This is the "real-time" aspect  
 What are the cost?  
   Does pointer reside in the from space? (If so, it needs to be copied)

- Requires test and branch for every dereference
- Difficult to run collector and program at same time
- Race condition between collector tracing heap and program threads
- Risk: two copies of the same object

Solution: let application use VM

<https://dl.acm.org/doi/10.1145/53990.53992>

Avoid explicit checks for references in application threads

- After copying roots, unmap the unscanned part of to-space

- Initially a page with roots etc.

- Page fault when thread accesses unscanned region

- handlers scans just the page and inspects all objects, then UNPROT

- at most one fault per page

- no compiler changes

Easy to make concurrent

- A collector thread can run concurrently with application thread

- A collector thread can UNPROT a page after scanning

- Only synchronization needed is for which thread is scanning which page

Are existing VM primitives good enough for concurrent GC?

- MAP2 is the only functionality issue -- but not really.

- We never have to make the same page accessible twice!

Are traps &c fast enough?

- They say no: 500 us to scan a page, 1200 us to take the trap.

- Why not scan 3 pages?

- How much slower to run Baker's actual algorithm, w/ checks?

- VM version might be faster! Even w/ slow traps.

- What about time saved by 2nd CPU scanning? They don't count this.

Is it an issue how often faults occur for concurrent GC?

- Not really -- more faults means more scanning.

- I.e. we'll get  $\leq$  one fault per page, at most.

Is user-level VM a good idea?

- Most of the use cases can be implemented by adding instructions

- E.g., check in application thread on reference

Pro:

- Avoid compiler changes

- CPU provides VM support

Con:

- Requires OS support, and efficient support

- Most OS kernel can't expose the raw hardware performance of paging

- OS imposes much abstraction

What has changed between 1991 and 2020?

- Tons of changes to VM API and implementations

- VM system evolves continuously

- Continued research too (e.g., see OSDI 2020)

Switching address space is free (tagged TLBs)

Extended addressability doesn't matter

- $2^{52}$  bytes of virtual address space